

Memory access patterns in Web Codecs

Current state and future developments

Paul Adenot

moz://a

W3C/SMPTE Joint Workshop on Professional Media
Production on the Web

memcpy is murder

Some numbers for a frame

- YUV420 1080p video frame SDR:
 $1920 * 1080 * 2 \approx \mathbf{4MB}$
- YUV420 4k video frame SDR:
 $3840 * 2160 * 2 \approx \mathbf{16MB}$
- P010 (\approx YUV420 10-bits) 4k video frame HDR:
 $3840 * 2160 * 4 \approx \mathbf{32MB}$

Time for a copy on really fast desktop workstation

AVX512 memcpy + DDR4 + optimized C++ = close to best case

- Hot caches
 - YUV420 1080p video frame SDR $\approx 1.5\text{ms}$
 - YUV420 4k video frame SDR $\approx 6.6\text{ms}$
 - P010 4k 10-bits video frame HDR $\approx 15\text{ms}$
- Cold caches
 - YUV420 1080p video frame SDR $\approx 4.5\text{ms}$
 - YUV420 4k video frame SDR $\approx 17\text{ms}$
 - P010 4k 10-bits video frame HDR $\approx 33\text{ms}$

GPU to CPU readback and upload

- Hardware decoded frames in GPU memory sometimes need to be copied to regular memory, this is very expensive
- It's always better to keep the VideoFrames on the GPU *if possible*

Why copy

- Custom post-processing in JavaScript or WASM
- Necessary to move data over to the WASM heap
- Sometimes necessary to work with other Web APIs

WebCodecs tries very hard to minimize copies

- Memory not explicitly visible: optimizations happen under the hood (e.g. copy on write), GPU surfaces are efficiently referenced
- Explicit copyTo methods to make it extra clear
- clone() method does *not* do a deep-copy

Necessary copies part 1 - easy fixes

- decode input: data is copied (WebCodecs issue #104)
- VideoFrame and AudioData copyTo: no way to "steal" the underlying memory yet (WebCodecs issue #287)
- memory cycling / allocator pressure (WebCodecs issue #212)

Buffer stealing

```
partial interface VideoFrame {  
  // closes the VideoFrame and transfer memory  
  Promise<ArrayBuffer> detach();  
};
```

(similar for `AudioData`).

Limit native allocator pressure

```
partial interface AudioDecoder {  
  // Detaches destination (need to be big enough)  
  // and write into it  
  // Detaches the memory in EncodedAudioChunk  
  undefined decode(EncodedAudioChunk chunk,  
    ArrayBuffer destination);  
};  
  
// `input` is the memory that was owned by `chunk`  
callback AudioDataOutputCallback =  
  undefined(AudioData output, ArrayBuffer input);
```

Necessary copies part 2 - harder problems

- Necessity to copy from/to the WASM heap
- Danger of SharedArrayBuffer vs. non-auditable codecs
- No read-only memory ranges
- No read-only memory: can't use memory ranges in encoder/decoders (BYOB)

Summary and positions at
<https://github.com/WICG/reducing-memory-copies>,
WebCodecs positions [issue #1](#).

Conclusion

There are problems, but there are also solutions in the works.

Generally, lots of common scenario work really well, but advance use-cases can be improved.